

5500 Final Project Group 3

```
In [1]: # Import necessary libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.preprocessing import LabelEncoder, StandardScaler
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score, confusion_matrix
from sklearn.utils import class_weight

# Deep Learning Libraries
import tensorflow as tf
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Dropout, Conv1D, MaxPooling1D, Flatten, LSTM, SimpleRNN
from tensorflow.keras.optimizers import Adam
from tensorflow.keras.callbacks import EarlyStopping
from tensorflow.keras.utils import to_categorical

# Set random seed for reproducibility
np.random.seed(42)
tf.random.set_seed(42)

# Load datasets (replace with your actual file paths)
try:
    train_data = pd.read_csv('trainset.csv')
    test_data = pd.read_csv('testset.csv')
except FileNotFoundError:
    print("Please make sure the files 'trainset.csv' and 'testset.csv' are in the correct directory.")
    raise

# Step 1: Data Exploration
print("\n=== Data Exploration ===\n")
print("Training set shape:", train_data.shape)
print("Test set shape:", test_data.shape)
print("\nTraining set class distribution:")
print(train_data['Subscribed'].value_counts())
print("\nTest set class distribution:")
```

```
print(test_data['Subscribed'].value_counts())

# Visualize class distribution
plt.figure(figsize=(10, 5))
plt.subplot(1, 2, 1)
train_data['Subscribed'].value_counts().plot(kind='bar', color=['skyblue', 'salmon'])
plt.title('Training Set Class Distribution')
plt.subplot(1, 2, 2)
test_data['Subscribed'].value_counts().plot(kind='bar', color=['skyblue', 'salmon'])
plt.title('Test Set Class Distribution')
plt.tight_layout()
plt.show()

# Explore categorical features
categorical_features = ['job', 'marital', 'education', 'housing', 'loan', 'contact', 'month', 'day_of_week', 'poutcom
plt.figure(figsize=(20, 25))
for i, feature in enumerate(categorical_features, 1):
    plt.subplot(5, 2, i)
    sns.countplot(x=feature, hue='Subscribed', data=train_data, palette='viridis')
    plt.title(f'{feature} Distribution')
    plt.xticks(rotation=45)
plt.tight_layout()
plt.show()

# Explore numerical features
numerical_features = ['age', 'duration', 'campaign', 'pdays', 'nr.employed']
plt.figure(figsize=(15, 10))
for i, feature in enumerate(numerical_features, 1):
    plt.subplot(2, 3, i)
    sns.boxplot(x='Subscribed', y=feature, data=train_data, palette='viridis')
    plt.title(f'{feature} Distribution')
plt.tight_layout()
plt.show()

# Step 2: Data Preprocessing
print("\n=== Data Preprocessing ===\n")

# Combine train and test for consistent preprocessing
combined = pd.concat([train_data, test_data], axis=0)

# Handle 'unknown' values - replace with mode or appropriate value
for column in categorical_features:
```

```
mode_val = combined[column].mode()[0]
combined[column] = combined[column].replace('unknown', mode_val)

# Convert pdays=999 to -1 (indicator for not previously contacted)
combined['pdays'] = combined['pdays'].replace(999, -1)

# Encode categorical variables
label_encoders = {}
for column in categorical_features:
    le = LabelEncoder()
    combined[column] = le.fit_transform(combined[column])
    label_encoders[column] = le

# Encode target variable
target_encoder = LabelEncoder()
combined['Subscribed'] = target_encoder.fit_transform(combined['Subscribed'])

# Split back into train and test
train_data = combined.iloc[:len(train_data)]
test_data = combined.iloc[len(train_data):]

# Separate features and target
X_train = train_data.drop('Subscribed', axis=1)
y_train = train_data['Subscribed']
X_test = test_data.drop('Subscribed', axis=1)
y_test = test_data['Subscribed']

# Scale numerical features
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

# Handle class imbalance
class_weights = class_weight.compute_class_weight('balanced', classes=np.unique(y_train), y=y_train)
class_weights = dict(enumerate(class_weights))

# Reshape data for CNN and RNN
X_train_resaped = X_train_scaled.reshape(X_train_scaled.shape[0], X_train_scaled.shape[1], 1)
X_test_resaped = X_test_scaled.reshape(X_test_scaled.shape[0], X_test_scaled.shape[1], 1)

# Step 3: Model Development and Training
print("\n=== Model Development and Training ===\n")
```

```
# MLP Model
def build_mlp():
    model = Sequential([
        Dense(64, activation='relu', input_shape=(X_train_scaled.shape[1],)),
        Dropout(0.3),
        Dense(32, activation='relu'),
        Dropout(0.2),
        Dense(1, activation='sigmoid')
    ])
    model.compile(optimizer=Adam(learning_rate=0.001),
                  loss='binary_crossentropy',
                  metrics=['accuracy'])
    return model

mlp_model = build_mlp()
print("MLP Model Summary:")
mlp_model.summary()

# CNN Model
def build_cnn():
    model = Sequential([
        Conv1D(filters=64, kernel_size=3, activation='relu', input_shape=(X_train_resaped.shape[1], 1)),
        MaxPooling1D(pool_size=2),
        Flatten(),
        Dense(32, activation='relu'),
        Dropout(0.2),
        Dense(1, activation='sigmoid')
    ])
    model.compile(optimizer=Adam(learning_rate=0.001),
                  loss='binary_crossentropy',
                  metrics=['accuracy'])
    return model

cnn_model = build_cnn()
print("\nCNN Model Summary:")
cnn_model.summary()

# LSTM Model
def build_lstm():
    model = Sequential([
        LSTM(64, input_shape=(X_train_resaped.shape[1], 1), return_sequences=True),
```

```
        Dropout(0.3),
        LSTM(32),
        Dropout(0.2),
        Dense(1, activation='sigmoid')
    ])
    model.compile(optimizer=Adam(learning_rate=0.001),
                  loss='binary_crossentropy',
                  metrics=['accuracy'])
    return model

lstm_model = build_lstm()
print("\nLSTM Model Summary:")
lstm_model.summary()

# RNN Model
def build_rnn():
    model = Sequential([
        SimpleRNN(64, input_shape=(X_train_resaped.shape[1], 1)),
        Dropout(0.3),
        Dense(32, activation='relu'),
        Dropout(0.2),
        Dense(1, activation='sigmoid')
    ])
    model.compile(optimizer=Adam(learning_rate=0.001),
                  loss='binary_crossentropy',
                  metrics=['accuracy'])
    return model

rnn_model = build_rnn()
print("\nRNN Model Summary:")
rnn_model.summary()

# Define early stopping callback
early_stopping = EarlyStopping(monitor='val_loss', patience=5, restore_best_weights=True)

# Train MLP
print("\nTraining MLP Model...")
mlp_history = mlp_model.fit(X_train_scaled, y_train,
                            validation_split=0.2,
                            epochs=50,
                            batch_size=64,
                            class_weight=class_weights,
```

```
        callbacks=[early_stopping],
        verbose=1)

# Train CNN
print("\nTraining CNN Model...")
cnn_history = cnn_model.fit(X_train_resaped, y_train,
                            validation_split=0.2,
                            epochs=50,
                            batch_size=64,
                            class_weight=class_weights,
                            callbacks=[early_stopping],
                            verbose=1)

# Train LSTM
print("\nTraining LSTM Model...")
lstm_history = lstm_model.fit(X_train_resaped, y_train,
                              validation_split=0.2,
                              epochs=50,
                              batch_size=64,
                              class_weight=class_weights,
                              callbacks=[early_stopping],
                              verbose=1)

# Train RNN
print("\nTraining RNN Model...")
rnn_history = rnn_model.fit(X_train_resaped, y_train,
                            validation_split=0.2,
                            epochs=50,
                            batch_size=64,
                            class_weight=class_weights,
                            callbacks=[early_stopping],
                            verbose=1)

# Step 4: Model Evaluation and Comparison
print("\n=== Model Evaluation and Comparison ===\n")

def evaluate_model(model, X_test, y_test, model_name):
    y_pred_prob = model.predict(X_test)
    y_pred = (y_pred_prob > 0.5).astype(int)

    accuracy = accuracy_score(y_test, y_pred)
    precision = precision_score(y_test, y_pred)
```

```
recall = recall_score(y_test, y_pred)
f1 = f1_score(y_test, y_pred)

print(f"{model_name} Performance:")
print(f"Accuracy: {accuracy:.4f}")
print(f"Precision: {precision:.4f}")
print(f"Recall: {recall:.4f}")
print(f"F1 Score: {f1:.4f}")

# Plot confusion matrix
cm = confusion_matrix(y_test, y_pred)
plt.figure(figsize=(6, 4))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues',
            xticklabels=['No', 'Yes'],
            yticklabels=['No', 'Yes'])
plt.title(f'{model_name} Confusion Matrix')
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.show()

    return accuracy, precision, recall, f1

# Evaluate MLP
mlp_metrics = evaluate_model(mlp_model, X_test_scaled, y_test, 'MLP')

# Evaluate CNN
cnn_metrics = evaluate_model(cnn_model, X_test_reshaped, y_test, 'CNN')

# Evaluate LSTM
lstm_metrics = evaluate_model(lstm_model, X_test_reshaped, y_test, 'LSTM')

# Evaluate RNN
rnn_metrics = evaluate_model(rnn_model, X_test_reshaped, y_test, 'RNN')

# Plot training history for all models
def plot_history(history, model_name):
    plt.figure(figsize=(12, 4))
    plt.subplot(1, 2, 1)
    plt.plot(history.history['accuracy'], label='Train Accuracy')
    plt.plot(history.history['val_accuracy'], label='Validation Accuracy')
    plt.title(f'{model_name} Accuracy')
    plt.xlabel('Epoch')
```

```
plt.ylabel('Accuracy')
plt.legend()

plt.subplot(1, 2, 2)
plt.plot(history.history['loss'], label='Train Loss')
plt.plot(history.history['val_loss'], label='Validation Loss')
plt.title(f'{model_name} Loss')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.legend()
plt.tight_layout()
plt.show()

plot_history(mlp_history, 'MLP')
plot_history(cnn_history, 'CNN')
plot_history(lstm_history, 'LSTM')
plot_history(rnn_history, 'RNN')

# Compare all models' performance
metrics_df = pd.DataFrame({
    'Model': ['MLP', 'CNN', 'LSTM', 'RNN'],
    'Accuracy': [mlp_metrics[0], cnn_metrics[0], lstm_metrics[0], rnn_metrics[0]],
    'Precision': [mlp_metrics[1], cnn_metrics[1], lstm_metrics[1], rnn_metrics[1]],
    'Recall': [mlp_metrics[2], cnn_metrics[2], lstm_metrics[2], rnn_metrics[2]],
    'F1 Score': [mlp_metrics[3], cnn_metrics[3], lstm_metrics[3], rnn_metrics[3]]
})

print("\n=== Model Performance Comparison ===")
print(metrics_df)

# Visual comparison of model metrics
plt.figure(figsize=(15, 8))
metrics = ['Accuracy', 'Precision', 'Recall', 'F1 Score']
for i, metric in enumerate(metrics, 1):
    plt.subplot(2, 2, i)
    sns.barplot(x='Model', y=metric, data=metrics_df)
    plt.title(metric)
    plt.ylim(0, 1)
plt.tight_layout()
plt.show()

# Step 5: Final Test Results
```



```
print("\n=== Final Test Results ===\n")
print("The models have been evaluated on the test set. The performance metrics are summarized above.")
print("The best performing model based on F1 Score is:", metrics_df.loc[metrics_df['F1 Score'].idxmax(), 'Model'])
```

=== Data Exploration ===

Training set shape: (29271, 15)

Test set shape: (11917, 15)

Training set class distribution:

Subscribed

no 26075

yes 3196

Name: count, dtype: int64

Test set class distribution:



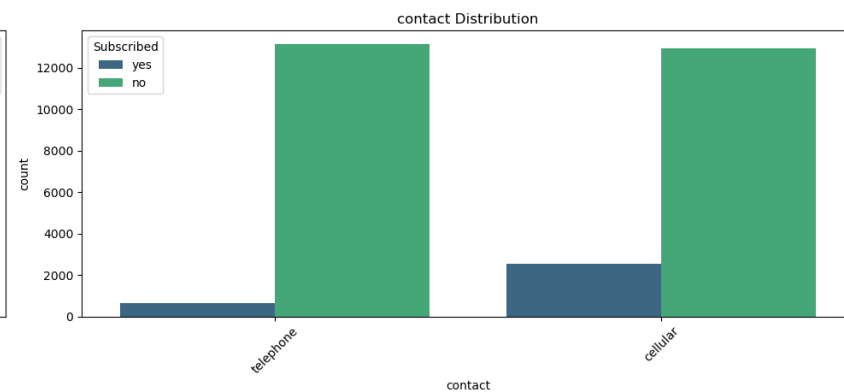





Subscribed

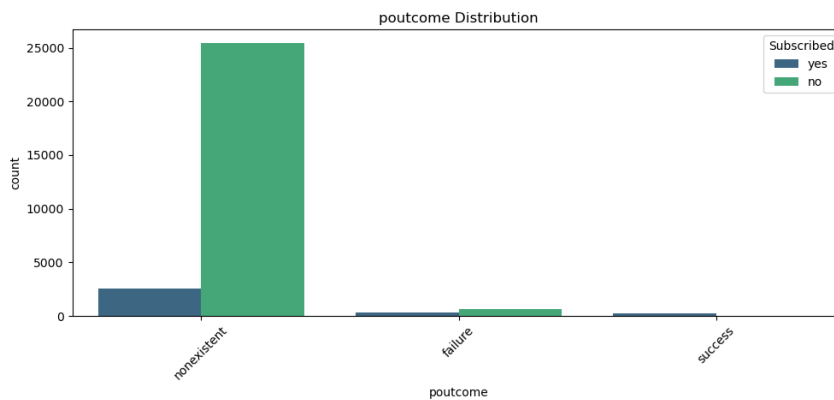


no 10473

yes 1444

Name: count, dtype: int64







C:\Users\Tom\AppData\Local\Temp\ipykernel_17180\2116035293.py:67: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.boxplot(x='Subscribed', y=feature, data=train_data, palette='viridis')
```

C:\Users\Tom\AppData\Local\Temp\ipykernel_17180\2116035293.py:67: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.boxplot(x='Subscribed', y=feature, data=train_data, palette='viridis')
```

C:\Users\Tom\AppData\Local\Temp\ipykernel_17180\2116035293.py:67: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.boxplot(x='Subscribed', y=feature, data=train_data, palette='viridis')
```

C:\Users\Tom\AppData\Local\Temp\ipykernel_17180\2116035293.py:67: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue`

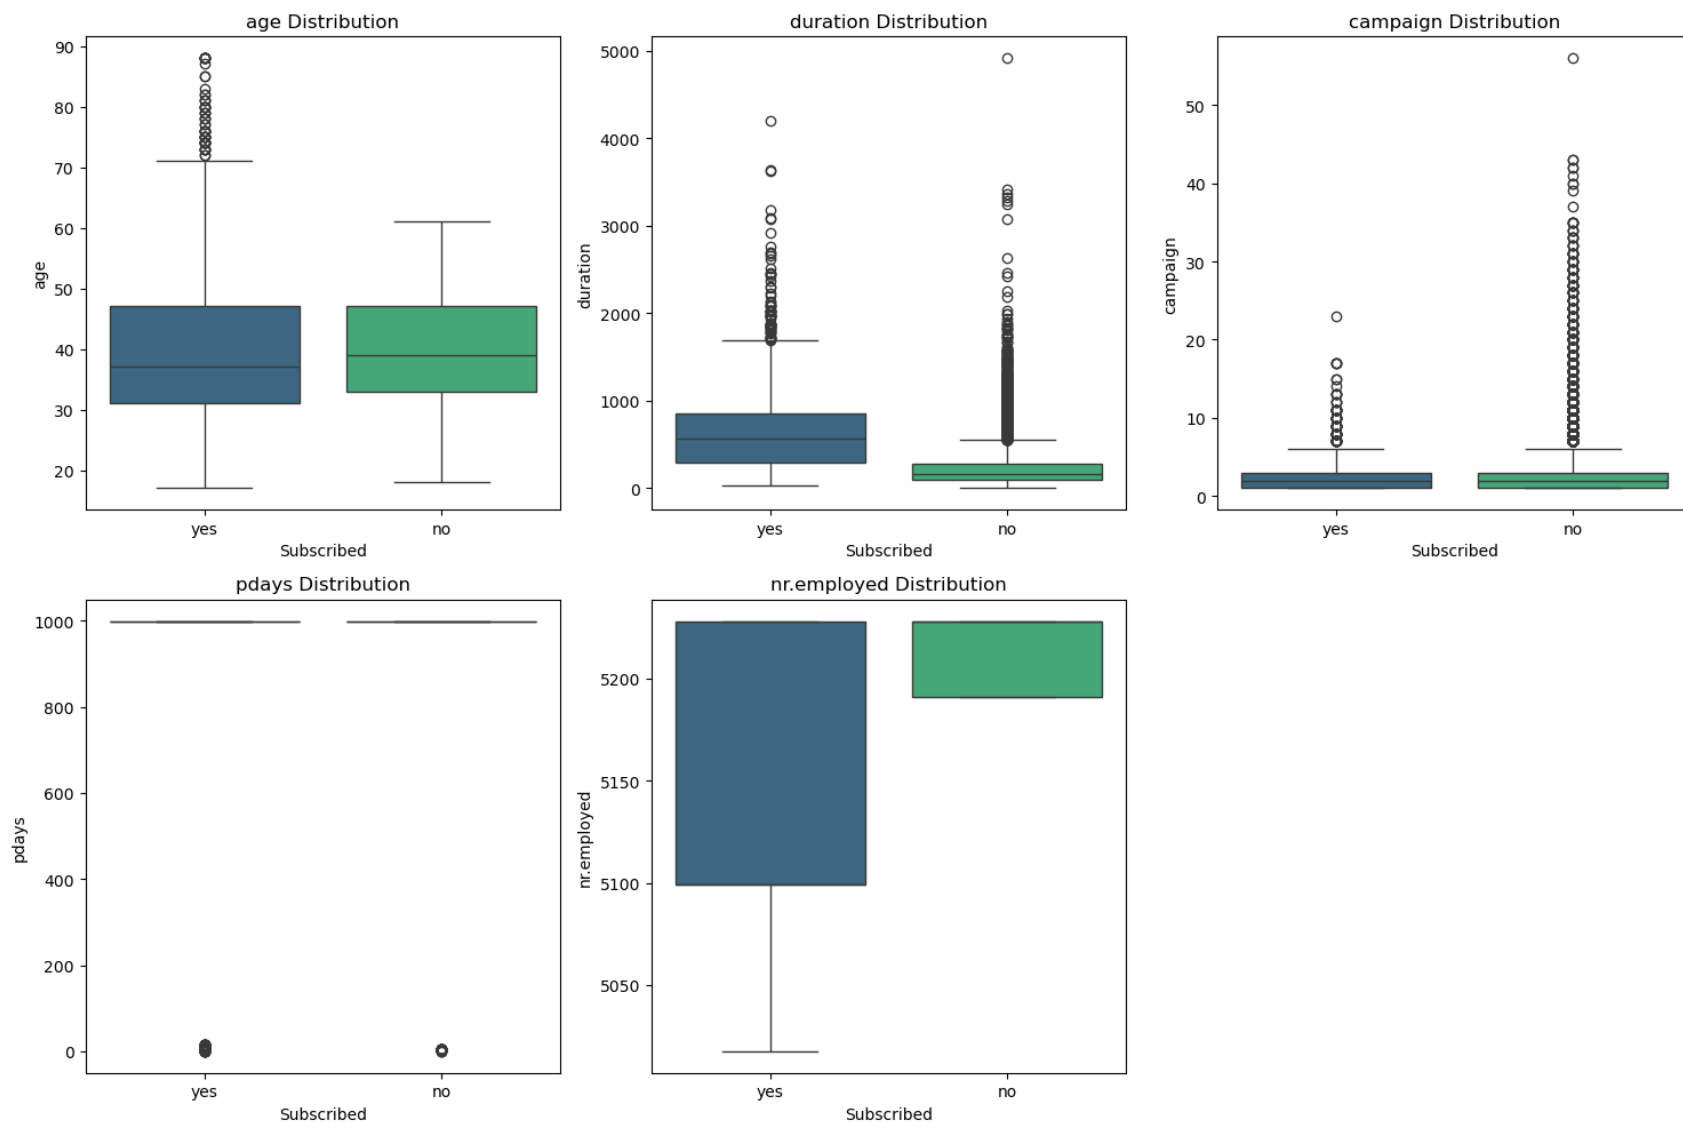
ue` and set `legend=False` for the same effect.

```
sns.boxplot(x='Subscribed', y=feature, data=train_data, palette='viridis')
```

C:\Users\Tom\AppData\Local\Temp\ipykernel_17180\2116035293.py:67: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.boxplot(x='Subscribed', y=feature, data=train_data, palette='viridis')
```



=== Data Preprocessing ===

=== Model Development and Training ===

```
C:\Users\Tom\anaconda3\Lib\site-packages\keras\src\layers\core\dense.py:87: UserWarning: Do not pass an `input_shape`  
`/`input_dim` argument to a layer. When using Sequential models, prefer using an `Input(shape)` object as the first l  
ayer in the model instead.
```

```
super().__init__(activity_regularizer=activity_regularizer, **kwargs)
```

MLP Model Summary:

Model: "sequential"

Layer (type)	Output Shape	Param #
dense (Dense)	(None, 64)	960
dropout (Dropout)	(None, 64)	0
dense_1 (Dense)	(None, 32)	2,080
dropout_1 (Dropout)	(None, 32)	0
dense_2 (Dense)	(None, 1)	33

Total params: 3,073 (12.00 KB)

Trainable params: 3,073 (12.00 KB)

Non-trainable params: 0 (0.00 B)

CNN Model Summary:

```
C:\Users\Tom\anaconda3\Lib\site-packages\keras\src\layers\convolutional\base_conv.py:107: UserWarning: Do not pass an  
`input_shape`/`input_dim` argument to a layer. When using Sequential models, prefer using an `Input(shape)` object as  
the first layer in the model instead.
```

```
super().__init__(activity_regularizer=activity_regularizer, **kwargs)
```

Model: "sequential_1"

Layer (type)	Output Shape	Param #
conv1d (Conv1D)	(None , 12, 64)	256
max_pooling1d (MaxPooling1D)	(None , 6, 64)	0
flatten (Flatten)	(None , 384)	0
dense_3 (Dense)	(None , 32)	12,320
dropout_2 (Dropout)	(None , 32)	0
dense_4 (Dense)	(None , 1)	33

Total params: 12,609 (49.25 KB)

Trainable params: 12,609 (49.25 KB)

Non-trainable params: 0 (0.00 B)

LSTM Model Summary:

C:\Users\Tom\anaconda3\Lib\site-packages\keras\src\layers\rnn\rnn.py:200: UserWarning: Do not pass an `input_shape`/`input_dim` argument to a layer. When using Sequential models, prefer using an `Input(shape)` object as the first layer in the model instead.

super().__init__(**kwargs)

Model: "sequential_2"

Layer (type)	Output Shape	Param #
lstm (LSTM)	(None , 14, 64)	16,896
dropout_3 (Dropout)	(None , 14, 64)	0
lstm_1 (LSTM)	(None , 32)	12,416
dropout_4 (Dropout)	(None , 32)	0
dense_5 (Dense)	(None , 1)	33

Total params: 29,345 (114.63 KB)

Trainable params: 29,345 (114.63 KB)

Non-trainable params: 0 (0.00 B)

RNN Model Summary:

Model: "sequential_3"

Layer (type)	Output Shape	Param #
simple_rnn (SimpleRNN)	(None, 64)	4,224
dropout_5 (Dropout)	(None, 64)	0
dense_6 (Dense)	(None, 32)	2,080
dropout_6 (Dropout)	(None, 32)	0
dense_7 (Dense)	(None, 1)	33


Total params: 6,337 (24.75 KB)

Trainable params: 6,337 (24.75 KB)


Non-trainable params: 0 (0.00 B)

Training MLP Model...


Epoch 1/50

366/366  **4s** 5ms/step - accuracy: 0.7181 - loss: 0.5234 - val_accuracy: 0.7522 - val_loss: 0.5797


Epoch 2/50

366/366  **1s** 4ms/step - accuracy: 0.9186 - loss: 0.2326 - val_accuracy: 0.7339 - val_loss: 0.6220


Epoch 3/50

366/366  **1s** 4ms/step - accuracy: 0.9196 - loss: 0.2105 - val_accuracy: 0.6499 - val_loss: 0.9073


Epoch 4/50

366/366  **2s** 4ms/step - accuracy: 0.9171 - loss: 0.1929 - val_accuracy: 0.5908 - val_loss: 1.1683

Epoch 5/50

366/366  **2s** 4ms/step - accuracy: 0.9178 - loss: 0.1831 - val_accuracy: 0.4941 - val_loss: 1.7673

Epoch 6/50


366/366  **2s** 4ms/step - accuracy: 0.9172 - loss: 0.1785 - val_accuracy: 0.4634 - val_loss: 2.6099

Training CNN Model...


Epoch 1/50

366/366  **4s** 6ms/step - accuracy: 0.7965 - loss: 0.3702 - val_accuracy: 0.8167 - val_loss: 1.0615


Epoch 2/50

366/366  **2s** 6ms/step - accuracy: 0.9227 - loss: 0.2183 - val_accuracy: 0.7667 - val_loss: 1.3201

Epoch 3/50

366/366  **2s** 6ms/step - accuracy: 0.9260 - loss: 0.1889 - val_accuracy: 0.5974 - val_loss: 1.8204

Epoch 4/50

366/366  **2s** 6ms/step - accuracy: 0.9254 - loss: 0.1769 - val_accuracy: 0.5132 - val_loss: 2.3651

Epoch 5/50


366/366  **2s** 7ms/step - accuracy: 0.9264 - loss: 0.1661 - val_accuracy: 0.4893 - val_loss: 3.0741

Epoch 6/50

366/366  **2s** 6ms/step - accuracy: 0.9273 - loss: 0.1587 - val_accuracy: 0.4811 - val_loss: 3.2162


Training LSTM Model...

Epoch 1/50


366/366  **13s** 23ms/step - accuracy: 0.4916 - loss: 0.5874 - val_accuracy: 0.6767 - val_loss: 0.979

2


Epoch 2/50

366/366  **9s** 24ms/step - accuracy: 0.8739 - loss: 0.2883 - val_accuracy: 0.6372 - val_loss: 1.1410


Epoch 3/50

366/366  **9s** 23ms/step - accuracy: 0.8942 - loss: 0.2622 - val_accuracy: 0.6835 - val_loss: 1.0677


Epoch 4/50

366/366  **9s** 24ms/step - accuracy: 0.9135 - loss: 0.2405 - val_accuracy: 0.6962 - val_loss: 1.0065

Epoch 5/50

366/366  **9s** 23ms/step - accuracy: 0.9171 - loss: 0.2283 - val_accuracy: 0.6989 - val_loss: 1.0192

Epoch 6/50

366/366  **8s** 23ms/step - accuracy: 0.9203 - loss: 0.2191 - val_accuracy: 0.6994 - val_loss: 1.0459

Training RNN Model...


Epoch 1/50

366/366  5s 8ms/step - accuracy: 0.8590 - loss: 0.3479 - val_accuracy: 0.7289 - val_loss: 0.5845


Epoch 2/50

366/366  3s 7ms/step - accuracy: 0.9136 - loss: 0.2147 - val_accuracy: 0.7392 - val_loss: 0.6096


Epoch 3/50

366/366  3s 7ms/step - accuracy: 0.9131 - loss: 0.2064 - val_accuracy: 0.6227 - val_loss: 0.8852


Epoch 4/50

366/366  3s 7ms/step - accuracy: 0.9132 - loss: 0.1894 - val_accuracy: 0.5474 - val_loss: 1.2928

Epoch 5/50

366/366  3s 7ms/step - accuracy: 0.9171 - loss: 0.1757 - val_accuracy: 0.5095 - val_loss: 1.8977

Epoch 6/50

366/366  3s 7ms/step - accuracy: 0.9172 - loss: 0.1708 - val_accuracy: 0.4883 - val_loss: 2.0044

=== Model Evaluation and Comparison ===

373/373  0s 1ms/step


MLP Performance:

Accuracy: 0.1385

Precision: 0.1233

Recall: 1.0000

F1 Score: 0.2196



373/373  1s 3ms/step


CNN Performance:

Accuracy: 0.1384

Precision: 0.1233

Recall: 1.0000

F1 Score: 0.2195



373/373 ————— 3s 8ms/step


LSTM Performance:

Accuracy: 0.1355

Precision: 0.1229

Recall: 1.0000

F1 Score: 0.2190



373/373 ————— 1s 3ms/step




RNN Performance:





Accuracy: 0.1354



Precision: 0.1229

Recall: 1.0000

F1 Score: 0.2189

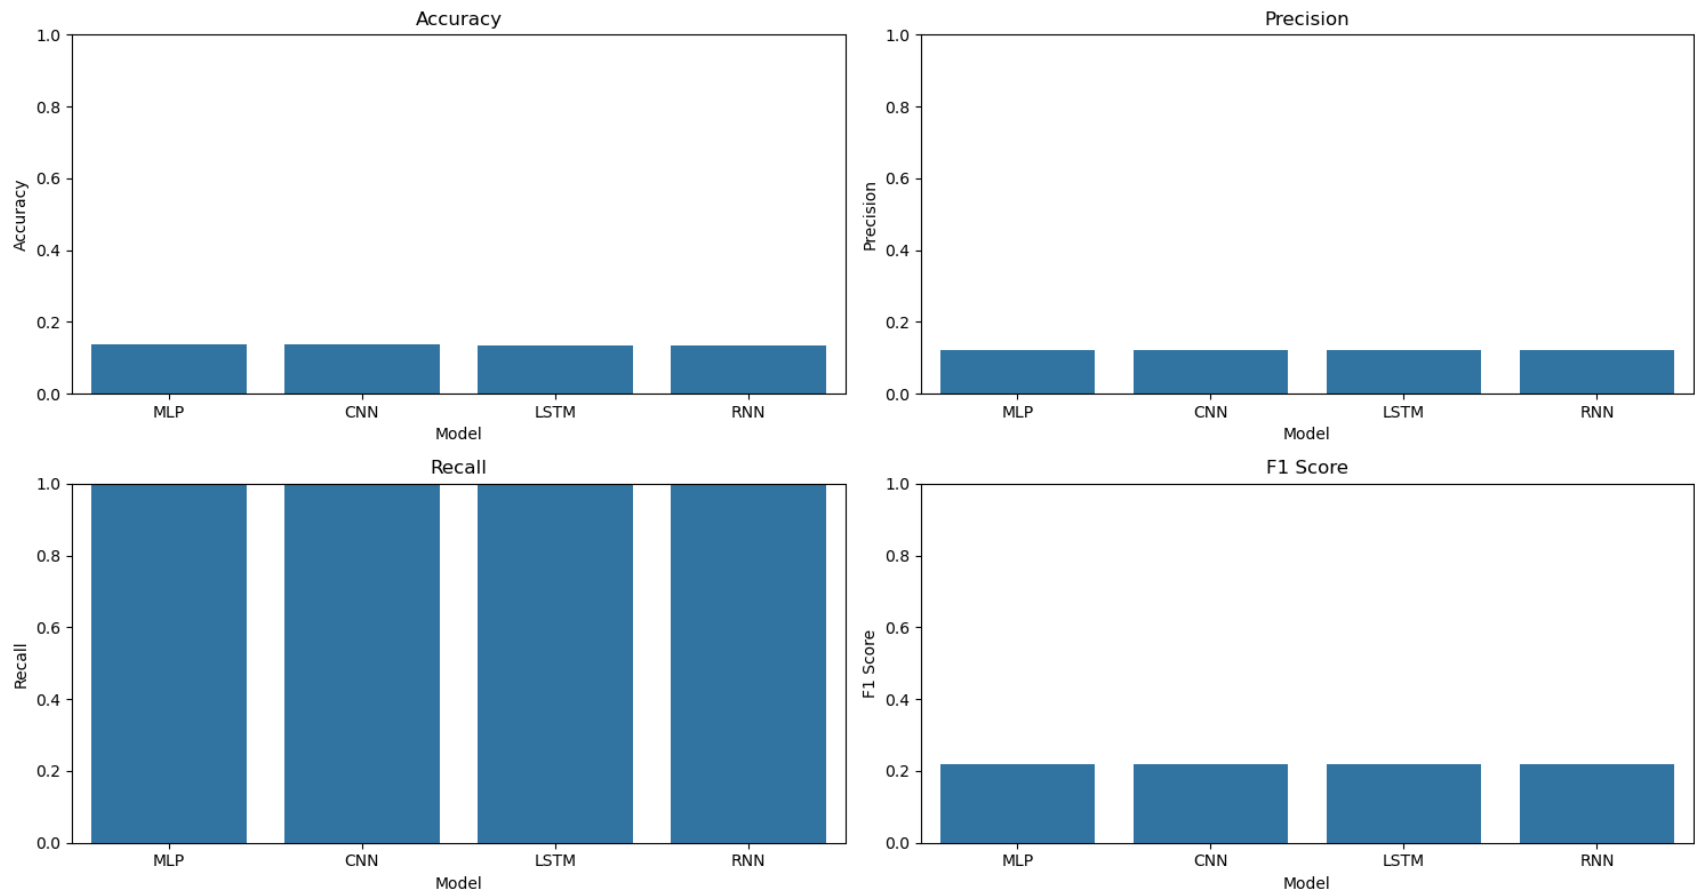






=== Model Performance Comparison ===

	Model	Accuracy	Precision	Recall	F1 Score
0	MLP	0.138542	0.123313	1.0	0.219553
1	CNN	0.138374	0.123292	1.0	0.219520
2	LSTM	0.135521	0.122935	1.0	0.218954
3	RNN	0.135353	0.122915	1.0	0.218921



=== Final Test Results ===

The models have been evaluated on the test set. The performance metrics are summarized above.
The best performing model based on F1 Score is: MLP

```
In [21]: # Calculate correlations for numeric features vs. target
numeric_data = data[numerical_features + ["Subscribed"]]
numeric_data["Subscribed"] = numeric_data["Subscribed"].map({"no": 0, "yes": 1}) # Encode target

correlation_matrix = numeric_data.corr(method="pearson")
target_correlations = correlation_matrix["Subscribed"].drop("Subscribed").sort_values(ascending=False)

# Plot correlation coefficients
plt.figure(figsize=(8, 4))
sns.barplot(x=target_correlations.values, y=target_correlations.index, palette="viridis")
```

```
plt.title("Correlation of Numeric Features with 'subscribed'")
plt.xlabel("Pearson Correlation Coefficient")
plt.ylabel("Feature")
plt.axvline(0, color="k", linestyle="--")
plt.show()
```

C:\Users\Tom\AppData\Local\Temp\ipykernel_28516\3013124865.py:3: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.

Try using `.loc[row_indexer,col_indexer] = value` instead


See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
numeric_data["Subscribed"] = numeric_data["Subscribed"].map({"no": 0, "yes": 1}) # Encode target
```

C:\Users\Tom\AppData\Local\Temp\ipykernel_28516\3013124865.py:10: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `y` variable to `hue` and set `legend=False` for the same effect.

```
sns.barplot(x=target_correlations.values, y=target_correlations.index, palette="viridis")
```



```
In [23]: # Target-encode categorical features for correlation-like analysis
categorical_data = data[categorical_features + ["Subscribed"]]
categorical_data["Subscribed"] = categorical_data["Subscribed"].map({"no": 0, "yes": 1})

# Calculate mean subscription rate per category (proxy for correlation)
mean_effects = {}
for col in categorical_features:
    mean_effects[col] = categorical_data.groupby(col)["Subscribed"].mean().sort_values(ascending=False)

# Plot top influential categories (e.g., for 'poutcome')
plt.figure(figsize=(10, 4))
mean_effects["poutcome"].plot(kind="bar", color="skyblue")
plt.title("Subscription Rate by 'poutcome' (Categorical Influence)")
plt.ylabel("Subscription Rate (Mean)")
plt.axhline(y=categorical_data["Subscribed"].mean(), color="r", linestyle="--", label="Overall Mean")
plt.legend()
plt.show()
```


C:\Users\Tom\AppData\Local\Temp\ipykernel_28516\901812262.py:3: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.


Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy


```
categorical_data["Subscribed"] = categorical_data["Subscribed"].map({"no": 0, "yes": 1})
```





```
In [25]: # Plot full correlation matrix (numeric features only)
plt.figure(figsize=(10, 6))
sns.heatmap(correlation_matrix, annot=True, cmap="coolwarm", vmin=-1, vmax=1, center=0)
plt.title("Correlation Matrix of Numeric Features")
plt.show()
```





```
In [27]: # Generate plots for all categorical variables
for col in categorical_features:
    plt.figure(figsize=(8, 3))
    mean_effects[col].plot(kind="bar", title=f"Subscription Rate by '{col}'")
    plt.axhline(y=categorical_data["Subscribed"].mean(), color="r", linestyle="--")
    plt.show()
```


















```
In [1]: # MLP, CNN and a Linear enhanced model for comparison

# Import necessary libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.preprocessing import LabelEncoder, StandardScaler
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score, confusion_matrix
from sklearn.utils import class_weight

# Deep Learning Libraries
import tensorflow as tf
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Dropout, Conv1D, MaxPooling1D, Flatten
from tensorflow.keras.optimizers import Adam
from tensorflow.keras.callbacks import EarlyStopping
from tensorflow.keras.utils import to_categorical

# Set random seed for reproducibility
np.random.seed(42)
tf.random.set_seed(42)

# Load datasets (replace with your actual file paths)
try:
    train_data = pd.read_csv('trainset.csv')
    test_data = pd.read_csv('testset.csv')
except FileNotFoundError:
    print("Please make sure the files 'trainset.csv' and 'testset.csv' are in the correct directory.")
    raise

# Step 1: Data Exploration
print("\n=== Data Exploration ===\n")
print("Training set shape:", train_data.shape)
print("Test set shape:", test_data.shape)
print("\nTraining set class distribution:")
print(train_data['Subscribed'].value_counts())
print("\nTest set class distribution:")
print(test_data['Subscribed'].value_counts())
```

```
# Visualize class distribution
plt.figure(figsize=(10, 5))
plt.subplot(1, 2, 1)
train_data['Subscribed'].value_counts().plot(kind='bar', color=['skyblue', 'salmon'])
plt.title('Training Set Class Distribution')
plt.subplot(1, 2, 2)
test_data['Subscribed'].value_counts().plot(kind='bar', color=['skyblue', 'salmon'])
plt.title('Test Set Class Distribution')
plt.tight_layout()
plt.show()

# Explore categorical features
categorical_features = ['job', 'marital', 'education', 'housing', 'loan', 'contact', 'month', 'day_of_week', 'poutcon
plt.figure(figsize=(20, 25))
for i, feature in enumerate(categorical_features, 1):
    plt.subplot(5, 2, i)
    sns.countplot(x=feature, hue='Subscribed', data=train_data, palette='viridis')
    plt.title(f'{feature} Distribution')
    plt.xticks(rotation=45)
plt.tight_layout()
plt.show()

# Explore numerical features
numerical_features = ['age', 'duration', 'campaign', 'pdays', 'nr.employed']
plt.figure(figsize=(15, 10))
for i, feature in enumerate(numerical_features, 1):
    plt.subplot(2, 3, i)
    sns.boxplot(x='Subscribed', y=feature, data=train_data, palette='viridis')
    plt.title(f'{feature} Distribution')
plt.tight_layout()
plt.show()

# Step 2: Data Preprocessing
print("\n=== Data Preprocessing ===\n")

# Combine train and test for consistent preprocessing
combined = pd.concat([train_data, test_data], axis=0)

# Handle 'unknown' values - replace with mode or appropriate value
for column in categorical_features:
    mode_val = combined[column].mode()[0]
    combined[column] = combined[column].replace('unknown', mode_val)
```

```
# Convert pdays=999 to -1 (indicator for not previously contacted)
combined['pdays'] = combined['pdays'].replace(999, -1)

# Encode categorical variables
label_encoders = {}
for column in categorical_features:
    le = LabelEncoder()
    combined[column] = le.fit_transform(combined[column])
    label_encoders[column] = le

# Encode target variable
target_encoder = LabelEncoder()
combined['Subscribed'] = target_encoder.fit_transform(combined['Subscribed'])

# Split back into train and test
train_data = combined.iloc[:len(train_data)]
test_data = combined.iloc[len(train_data):]

# Separate features and target
X_train = train_data.drop('Subscribed', axis=1)
y_train = train_data['Subscribed']
X_test = test_data.drop('Subscribed', axis=1)
y_test = test_data['Subscribed']

# Scale numerical features
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

# Handle class imbalance
class_weights = class_weight.compute_class_weight('balanced', classes=np.unique(y_train), y=y_train)
class_weights = dict(enumerate(class_weights))

# Reshape data for CNN
X_train_resaped = X_train_scaled.reshape(X_train_scaled.shape[0], X_train_scaled.shape[1], 1)
X_test_resaped = X_test_scaled.reshape(X_test_scaled.shape[0], X_test_scaled.shape[1], 1)

# Step 3: Model Development and Training
print("\n=== Model Development and Training ===\n")

# MLP Model
```



```
def build_mlp():
    model = Sequential([
        Dense(64, activation='relu', input_shape=(X_train_scaled.shape[1],)),
        Dropout(0.3),
        Dense(32, activation='relu'),
        Dropout(0.2),
        Dense(1, activation='sigmoid')
    ])
    model.compile(optimizer=Adam(learning_rate=0.001),
                  loss='binary_crossentropy',
                  metrics=['accuracy'])
    return model

mlp_model = build_mlp()
print("MLP Model Summary:")
mlp_model.summary()

# CNN Model
def build_cnn():
    model = Sequential([
        Conv1D(filters=64, kernel_size=3, activation='relu', input_shape=(X_train_resaped.shape[1], 1)),
        MaxPooling1D(pool_size=2),
        Flatten(),
        Dense(32, activation='relu'),
        Dropout(0.2),
        Dense(1, activation='sigmoid')
    ])
    model.compile(optimizer=Adam(learning_rate=0.001),
                  loss='binary_crossentropy',
                  metrics=['accuracy'])
    return model

cnn_model = build_cnn()
print("\nCNN Model Summary:")
cnn_model.summary()

# Linear Model (a simple model that's generally not good for this type of data)
def build_linear():
    model = Sequential([
        Dense(1, activation='sigmoid', input_shape=(X_train_scaled.shape[1],))
    ])
    model.compile(optimizer=Adam(learning_rate=0.001),
```

```
        loss='binary_crossentropy',
        metrics=['accuracy'])
    return model

linear_model = build_linear()
print("\nLinear Model Summary:")
linear_model.summary()

# Define early stopping callback
early_stopping = EarlyStopping(monitor='val_loss', patience=5, restore_best_weights=True)

# Train MLP
print("\nTraining MLP Model...")
mlp_history = mlp_model.fit(X_train_scaled, y_train,
                           validation_split=0.2,
                           epochs=50,
                           batch_size=64,
                           class_weight=class_weights,
                           callbacks=[early_stopping],
                           verbose=1)

# Train CNN
print("\nTraining CNN Model...")
cnn_history = cnn_model.fit(X_train_rescaled, y_train,
                           validation_split=0.2,
                           epochs=50,
                           batch_size=64,
                           class_weight=class_weights,
                           callbacks=[early_stopping],
                           verbose=1)

# Train Linear
print("\nTraining Linear Model...")
linear_history = linear_model.fit(X_train_scaled, y_train,
                                 validation_split=0.2,
                                 epochs=50,
                                 batch_size=64,
                                 class_weight=class_weights,
                                 callbacks=[early_stopping],
                                 verbose=1)

# Step 4: Model Evaluation and Comparison
```

```
print("\n=== Model Evaluation and Comparison ===\n")

def evaluate_model(model, X_test, y_test, model_name):
    y_pred_prob = model.predict(X_test)
    y_pred = (y_pred_prob > 0.5).astype(int)

    accuracy = accuracy_score(y_test, y_pred)
    precision = precision_score(y_test, y_pred)
    recall = recall_score(y_test, y_pred)
    f1 = f1_score(y_test, y_pred)

    print(f"{model_name} Performance:")
    print(f"Accuracy: {accuracy:.4f}")
    print(f"Precision: {precision:.4f}")
    print(f"Recall: {recall:.4f}")
    print(f"F1 Score: {f1:.4f}")

    # Plot confusion matrix
    cm = confusion_matrix(y_test, y_pred)
    plt.figure(figsize=(6, 4))
    sns.heatmap(cm, annot=True, fmt='d', cmap='Blues',
                xticklabels=['No', 'Yes'],
                yticklabels=['No', 'Yes'])
    plt.title(f'{model_name} Confusion Matrix')
    plt.xlabel('Predicted')
    plt.ylabel('Actual')
    plt.show()

    return accuracy, precision, recall, f1

# Evaluate MLP
mlp_metrics = evaluate_model(mlp_model, X_test_scaled, y_test, 'MLP')

# Evaluate CNN
cnn_metrics = evaluate_model(cnn_model, X_test_rescaled, y_test, 'CNN')

# Evaluate Linear
linear_metrics = evaluate_model(linear_model, X_test_scaled, y_test, 'Linear')

# Plot training history for all models
def plot_history(history, model_name):
    plt.figure(figsize=(12, 4))
```

```
plt.subplot(1, 2, 1)
plt.plot(history.history['accuracy'], label='Train Accuracy')
plt.plot(history.history['val_accuracy'], label='Validation Accuracy')
plt.title(f'{model_name} Accuracy')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.legend()

plt.subplot(1, 2, 2)
plt.plot(history.history['loss'], label='Train Loss')
plt.plot(history.history['val_loss'], label='Validation Loss')
plt.title(f'{model_name} Loss')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.legend()
plt.tight_layout()
plt.show()

plot_history(mlp_history, 'MLP')
plot_history(cnn_history, 'CNN')
plot_history(linear_history, 'Linear')

# Compare all models' performance
metrics_df = pd.DataFrame({
    'Model': ['MLP', 'CNN', 'Linear'],
    'Accuracy': [mlp_metrics[0], cnn_metrics[0], linear_metrics[0]],
    'Precision': [mlp_metrics[1], cnn_metrics[1], linear_metrics[1]],
    'Recall': [mlp_metrics[2], cnn_metrics[2], linear_metrics[2]],
    'F1 Score': [mlp_metrics[3], cnn_metrics[3], linear_metrics[3]]
})

print("\n=== Model Performance Comparison ===")
print(metrics_df)

# Visual comparison of model metrics
plt.figure(figsize=(15, 8))
metrics = ['Accuracy', 'Precision', 'Recall', 'F1 Score']
for i, metric in enumerate(metrics, 1):
    plt.subplot(2, 2, i)
    sns.barplot(x='Model', y=metric, data=metrics_df)
    plt.title(metric)
    plt.ylim(0, 1)
```

```
plt.tight_layout()
plt.show()

# Step 5: Final Test Results
print("\n=== Final Test Results ===\n")
print("The models have been evaluated on the test set. The performance metrics are summarized above.")
print("The best performing model based on F1 Score is:", metrics_df.loc[metrics_df['F1 Score'].idxmax(), 'Model'])
```

=== Data Exploration ===

Training set shape: (29271, 15)

Test set shape: (11917, 15)

Training set class distribution:

Subscribed

no 26075

yes 3196

Name: count, dtype: int64

Test set class distribution:









Subscribed




no 10473

yes 1444

Name: count, dtype: int64







C:\Users\Tom\AppData\Local\Temp\ipykernel_11852\2599279041.py:69: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.boxplot(x='Subscribed', y=feature, data=train_data, palette='viridis')
```

C:\Users\Tom\AppData\Local\Temp\ipykernel_11852\2599279041.py:69: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.boxplot(x='Subscribed', y=feature, data=train_data, palette='viridis')
```

C:\Users\Tom\AppData\Local\Temp\ipykernel_11852\2599279041.py:69: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.boxplot(x='Subscribed', y=feature, data=train_data, palette='viridis')
```

C:\Users\Tom\AppData\Local\Temp\ipykernel_11852\2599279041.py:69: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue`

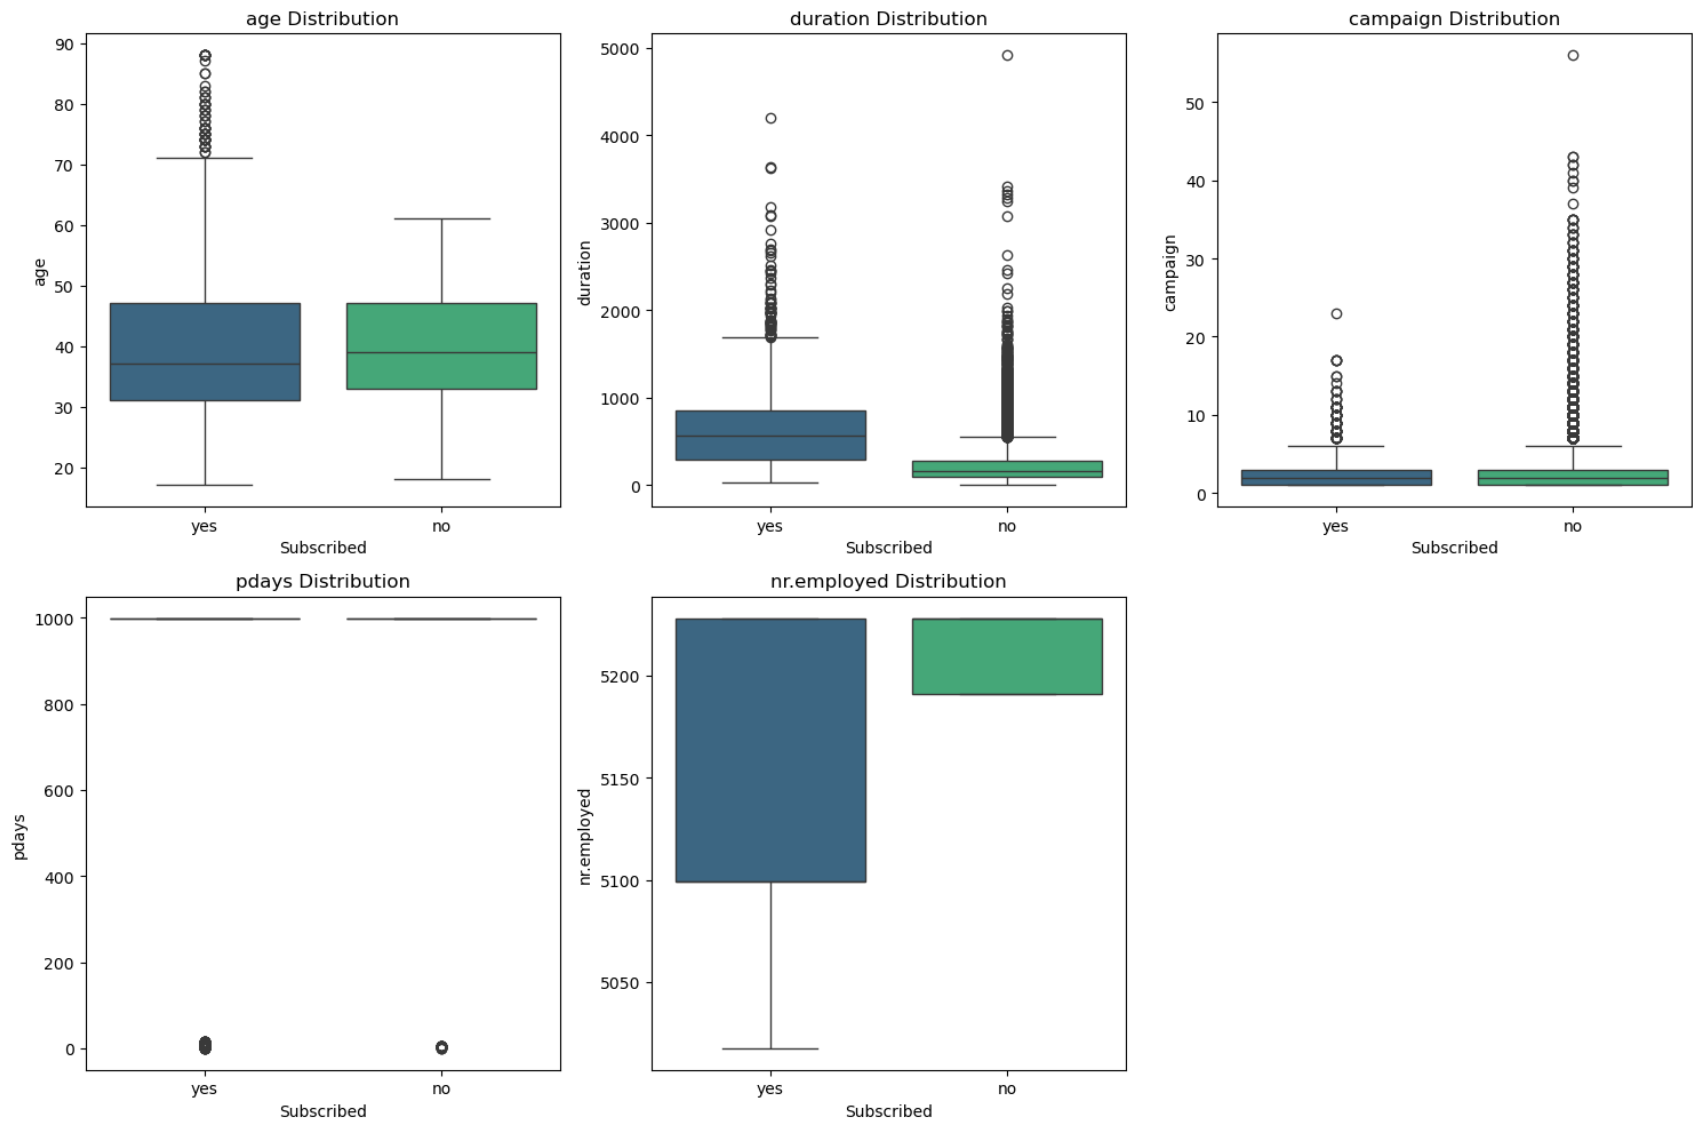
ue` and set `legend=False` for the same effect.

```
sns.boxplot(x='Subscribed', y=feature, data=train_data, palette='viridis')
```

C:\Users\Tom\AppData\Local\Temp\ipykernel_11852\2599279041.py:69: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.boxplot(x='Subscribed', y=feature, data=train_data, palette='viridis')
```



=== Data Preprocessing ===

=== Model Development and Training ===

C:\Users\Tom\anaconda3\Lib\site-packages\keras\src\layers\core\dense.py:87: UserWarning: Do not pass an `input_shape`/`input_dim` argument to a layer. When using Sequential models, prefer using an `Input(shape)` object as the first layer in the model instead.

```
super().__init__(activity_regularizer=activity_regularizer, **kwargs)
```

MLP Model Summary:

Model: "sequential"

Layer (type)	Output Shape	Param #
dense (Dense)	(None, 64)	960
dropout (Dropout)	(None, 64)	0
dense_1 (Dense)	(None, 32)	2,080
dropout_1 (Dropout)	(None, 32)	0
dense_2 (Dense)	(None, 1)	33

Total params: 3,073 (12.00 KB)

Trainable params: 3,073 (12.00 KB)

Non-trainable params: 0 (0.00 B)

CNN Model Summary:

C:\Users\Tom\anaconda3\Lib\site-packages\keras\src\layers\convolutional\base_conv.py:107: UserWarning: Do not pass an `input_shape`/`input_dim` argument to a layer. When using Sequential models, prefer using an `Input(shape)` object as the first layer in the model instead.

```
super().__init__(activity_regularizer=activity_regularizer, **kwargs)
```

Model: "sequential_1"

Layer (type)	Output Shape	Param #
conv1d (Conv1D)	(None , 12, 64)	256
max_pooling1d (MaxPooling1D)	(None , 6, 64)	0
flatten (Flatten)	(None , 384)	0
dense_3 (Dense)	(None , 32)	12,320
dropout_2 (Dropout)	(None , 32)	0
dense_4 (Dense)	(None , 1)	33

Total params: 12,609 (49.25 KB)

Trainable params: 12,609 (49.25 KB)

Non-trainable params: 0 (0.00 B)

Linear Model Summary:

Model: "sequential_2"

Layer (type)	Output Shape	Param #
dense_5 (Dense)	(None , 1)	15

Total params: 15 (60.00 B)

Trainable params: 15 (60.00 B)


Non-trainable params: 0 (0.00 B)

Training MLP Model...


Epoch 1/50

366/366  4s 4ms/step - accuracy: 0.7760 - loss: 0.4491 - val_accuracy: 0.7689 - val_loss: 0.5842


Epoch 2/50

366/366  2s 5ms/step - accuracy: 0.9131 - loss: 0.2270 - val_accuracy: 0.7209 - val_loss: 0.8300


Epoch 3/50

366/366  1s 4ms/step - accuracy: 0.9133 - loss: 0.2022 - val_accuracy: 0.6555 - val_loss: 1.1295


Epoch 4/50

366/366  1s 3ms/step - accuracy: 0.9162 - loss: 0.1870 - val_accuracy: 0.5221 - val_loss: 1.6430

Epoch 5/50


366/366  1s 4ms/step - accuracy: 0.9212 - loss: 0.1763 - val_accuracy: 0.4639 - val_loss: 2.6601

Epoch 6/50

366/366  1s 3ms/step - accuracy: 0.9221 - loss: 0.1710 - val_accuracy: 0.4576 - val_loss: 3.5997

Training CNN Model...

Epoch 1/50

366/366  4s 6ms/step - accuracy: 0.8303 - loss: 0.3637 - val_accuracy: 0.8029 - val_loss: 1.0018


Epoch 2/50

366/366  2s 5ms/step - accuracy: 0.9247 - loss: 0.2087 - val_accuracy: 0.7740 - val_loss: 1.2481


Epoch 3/50

366/366  2s 5ms/step - accuracy: 0.9247 - loss: 0.1915 - val_accuracy: 0.6319 - val_loss: 1.5834


Epoch 4/50

366/366  2s 6ms/step - accuracy: 0.9230 - loss: 0.1773 - val_accuracy: 0.5255 - val_loss: 1.9201

Epoch 5/50


366/366  2s 5ms/step - accuracy: 0.9241 - loss: 0.1718 - val_accuracy: 0.4864 - val_loss: 2.2493

Epoch 6/50


366/366  2s 5ms/step - accuracy: 0.9235 - loss: 0.1648 - val_accuracy: 0.4675 - val_loss: 2.5130

Training Linear Model...


Epoch 1/50

366/366  2s 3ms/step - accuracy: 0.6206 - loss: 0.7897 - val_accuracy: 0.1851 - val_loss: 1.0380

Epoch 2/50

366/366  1s 3ms/step - accuracy: 0.6812 - loss: 0.4970 - val_accuracy: 0.2789 - val_loss: 0.8548


Epoch 3/50

366/366  1s 3ms/step - accuracy: 0.7867 - loss: 0.3875 - val_accuracy: 0.5658 - val_loss: 0.7259


Epoch 4/50

366/366  1s 2ms/step - accuracy: 0.8897 - loss: 0.3319 - val_accuracy: 0.6941 - val_loss: 0.6467

Epoch 5/50

366/366  1s 3ms/step - accuracy: 0.9122 - loss: 0.2986 - val_accuracy: 0.7534 - val_loss: 0.5945

Epoch 6/50

366/366  1s 3ms/step - accuracy: 0.9210 - loss: 0.2769 - val_accuracy: 0.7855 - val_loss: 0.5578

Epoch 7/50

366/366  1s 3ms/step - accuracy: 0.9253 - loss: 0.2623 - val_accuracy: 0.8046 - val_loss: 0.5311
Epoch 8/50

366/366  1s 3ms/step - accuracy: 0.9282 - loss: 0.2521 - val_accuracy: 0.8162 - val_loss: 0.5114
Epoch 9/50

366/366  1s 3ms/step - accuracy: 0.9286 - loss: 0.2450 - val_accuracy: 0.8256 - val_loss: 0.4968
Epoch 10/50

366/366  1s 4ms/step - accuracy: 0.9297 - loss: 0.2399 - val_accuracy: 0.8292 - val_loss: 0.4862
Epoch 11/50

366/366  1s 4ms/step - accuracy: 0.9301 - loss: 0.2364 - val_accuracy: 0.8331 - val_loss: 0.4786
Epoch 12/50

366/366  1s 3ms/step - accuracy: 0.9305 - loss: 0.2339 - val_accuracy: 0.8359 - val_loss: 0.4733
Epoch 13/50

366/366  1s 3ms/step - accuracy: 0.9308 - loss: 0.2321 - val_accuracy: 0.8374 - val_loss: 0.4700
Epoch 14/50

366/366  1s 4ms/step - accuracy: 0.9311 - loss: 0.2308 - val_accuracy: 0.8372 - val_loss: 0.4680
Epoch 15/50

366/366  1s 3ms/step - accuracy: 0.9313 - loss: 0.2299 - val_accuracy: 0.8362 - val_loss: 0.4673
Epoch 16/50

366/366  1s 3ms/step - accuracy: 0.9313 - loss: 0.2293 - val_accuracy: 0.8360 - val_loss: 0.4676
Epoch 17/50


366/366  1s 4ms/step - accuracy: 0.9312 - loss: 0.2289 - val_accuracy: 0.8331 - val_loss: 0.4686
Epoch 18/50


366/366  1s 4ms/step - accuracy: 0.9311 - loss: 0.2285 - val_accuracy: 0.8307 - val_loss: 0.4703
Epoch 19/50

366/366  1s 4ms/step - accuracy: 0.9311 - loss: 0.2283 - val_accuracy: 0.8261 - val_loss: 0.4727
Epoch 20/50

366/366  1s 3ms/step - accuracy: 0.9311 - loss: 0.2281 - val_accuracy: 0.8210 - val_loss: 0.4756

=== Model Evaluation and Comparison ===

373/373  1s 2ms/step
MLP Performance:
Accuracy: 0.1362
Precision: 0.1230
Recall: 1.0000
F1 Score: 0.2191



373/373 1s 2ms/step


CNN Performance:

Accuracy: 0.1382

Precision: 0.1233

Recall: 1.0000

F1 Score: 0.2195



373/373 1s 2ms/step




Linear Performance:





Accuracy: 0.1386

Precision: 0.1233

Recall: 1.0000





F1 Score: 0.2196





=== Model Performance Comparison ===

	Model	Accuracy	Precision	Recall	F1 Score
0	MLP	0.136192	0.123019	1.0	0.219087
1	CNN	0.138206	0.123271	1.0	0.219486
2	Linear	0.138625	0.123324	1.0	0.219570



=== Final Test Results ===

The models have been evaluated on the test set. The performance metrics are summarized above.
The best performing model based on F1 Score is: Linear

In []:

In []: